




How We Protect Cat Memes from DDoS



Spencer Koch
Principal Security Engineer



Pratik Lotia
Senior Security Engineer



Context Time



What You'll Get From This Talk

- How Reddit thinks about ratelimiting architecture
- Signals and their importance
- Where to do ratelimiting
- Other resiliency strategies
- Lessons learned from 6 years of DDoS fights



Reddit's Kind of a Big Deal...



Per week: 1.3 Trillion requests, 175 Petabytes (83% cached)

Service	RPS
GraphQL	275k
Ads Pixel	100k+
Server Side Rendering	150k
Media	1200k



Where the DDoS Crew Came From...



2019 - evolved from a couple of OG engineers duct taping things together



2021 - our first Transport team, getting more serious about defenses

2025 - now both Transport and Traffic teams; dedicated engineers focused on DDoS mitigations and response





Signals



Signals - Foundational

- **Foundational:**
 - GeolP+
 - User Agent
 - Path, Query Params
- **TLS Fingerprints:**
 - JA3 / JA4 Calculations
 - GREASE - incompatibility vs randomness
 - Protocol sprawl



| <https://browserleaks.com/tls> - for more info

| <https://engineering.salesforce.com/tls-fingerprinting-with-ja3-and-ja3s-247362855967/> - JA3 details

| https://github.com/FoxIO-LLC/ja4/blob/main/technical_details/JA4.md - JA4 details



JA4: TLS Client Fingerprint

- Protocol: TCP = "t", QUIC = "q"
- TLS version: 1.2= "12", 1.3 = "13"
- SNI: SNI present = "d" (to domain), no SNI = "i" (to IP)
- Number of cipher suites
- Number of extensions
- First ALPN value (00 if no ALPN)

JA4: t13d1516h2_8daaf6152771_02713d6af862

JA4_a

JA4_b


JA4_c

- Truncated SHA256 hash of the cipher suites, sorted
- Truncated SHA256 hash of the extensions, sorted + signature algorithms

<https://blog.cloudflare.com/ja4-signals/>



Signals - Intermediate



- **Request Header Fingerprints**
 - Order, casing, and presence of headers
 - Must be done at receiving proxy before munging
 - Fastly - *fastly_info.oh_fingerprint*
 - Cloudflare - *request.cf?.requestHeaderNames*

| https://privacycheck.sec.lrz.de/passive/fp_hs/fp_header_signature.php - RHS info



Signals - Intermediate

```
Request Example (Chrome)
```

```
=====
Host: reddit.local
X-Forwarded-For: 192.168.56.1
Connection: close
Content-Length: 295
Accept: application/json, text/javascript, */*; q=0.01
Origin: https://reddit.local
X-Requested-With: XMLHttpRequest
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_11_2) AppleWebKit/537.36 (KHTML, like Ge
Content-Type: application/x-www-form-urlencoded; charset=UTF-8
Referer: https://reddit.local/new/
Accept-Encoding: gzip, deflate
Accept-Language: en-US,en;q=0.8
Cookie: loid=aj0uPPBsarNtdho6iK; loidcreated=2016-03-03T05%3A09%3A08.565Z; reddit_session=53%2
DNT: 1
X-Forwarded-Proto: https
```




```
request_header_signature': 'Am:GX:K6:Ac:AT:M9:Gj:A6:Af:Az:AQ:AR:Ag:Kv:CB
```



Signals - Advanced

- **Behavioral Fingerprinting**
 - Network request timings
 - Analytical events fired or missing
 - Cookies, JWT presence and validation
 - Login flow or “expected” order of operations
 - CAPTCHA or “someone else’s model”





Where to Ratelimit



Where to Ratelimit?

- Leave Layer 3 / 4 attacks to internet / cloud providers
- Difference between Edge and Application layer
 - **Edge:**
 - Cheaper, but lacks context of the app
 - User session, login state, limited business logic
 - **Application:**
 - More sophisticated logic and enriched data
 - Expensive as your compute does the work



Where to Ratelimit?

- Both need a KV store
 - Do you need more control over behavior?
 - Difficulty of flushing / unsetting?
 - Programmatic interactions?
- How to handle eventual consistency?
 - Consistent routing based on geo / session / stickiness
 - Circular hash ring algorithms




| <https://dev.to/mochafreddo/nginx-generic-hash-load-balancing-a-comprehensive-guide-3bca> - nginx hashing primer

Where to Ratelimit?

- Routing architecture principles:
 - “Reasonable” domain carve outs: media vs app split, chat vs main site, modmail
 - Cache where you can
 - Isolate workloads with specific purpose
- Gotta use both!







**Edge
Rate limiting**

Edge Ratelimiting

- **Fastly's Edge Rate Limiting**
 - Two KV stores:
 - **penaltybox**: a TTL for naughty elements
 - **ratecounter**: counter for observed hits
 - You determine what element to track (IP, TLS, etc.)
 - Edge dictionaries to control configurables (request rate, exclusions)



```
sub ip_rate_limit_run {  
    declare local var.ip_ratelimit_exceeded BOOL;  
    declare local var.window INTEGER;  
    declare local var.window_limit INTEGER;  
    declare local var.increment INTEGER;  
    declare local var.key STRING;  
  
    # Defaults (for comparison) with the TLS rate limiter.  
    set var.key = req.http.CDN-Client-IP;  
    set var.increment = 1;  
    set var.window = 10;  
    set var.window_limit = 15;  
  
    set var.ip_ratelimit_exceeded = ratelimit.check_rate(  
        var.key,  
        ip_counter_10s,  
        var.increment,      # Increment.  
        var.window,         # Sample window (in seconds)  
        var.window_limit,   # Limit in RPS  
        ip_pbox,  
        2m);
```

Example Fastly Construction



Edge Ratelimiting

- **Cloudflare**
 - WAF ratelimit rules - precanned and Enterprise only
 - Workers - use a KV and do it yourself
 - Don't forget your logging
 - `x-reddit-synthetic-code` & `x-reddit-synthetic-reason`
 - 429 is extremely terse signal to engineers



| <https://developers.cloudflare.com/waf/rate-limiting-rules/> - WAF rules




Edge Ratelimiting

- **Series of band-pass filters** with decreasing rate values:
IPs→TLS→pools→recaptcha
- **Weaknesses at Edge:**
 - CGNAT, shared IP blocks, TLS by user agent
 - Requires tuning - approved crawlers, third parties
 - Country+ASN reputation scoring or crowd-level heuristics? Depends on your business model
 - Collateral damage will happen! Revert fast

|







Application RateLimiting



Application Ratelimiting


- FIRST: **Common library for ratelimiting**
 - Pick an algorithm and provide guidance to your devs!
 - Sliding window algorithm with Redis TTL
 - Redis - common, cheap, make it easy for devs to use
 - “Allowance” and “interval” selection:
(how fast could a human do it) x (some reasonable multiplier)


| <https://github.com/reddit/baseplate.py/blob/develop/baseplate/lib/ratelimit/ratelimit.py> - Reddit's implementation



Application Ratelimiting

- Best for:
 - “Per user, per endpoint” logic
 - OAuth ratelimits per client ID
- API route handling (front) vs datastore execution (back)?
- Square wave signals for abuse detection
- Worldwide Reddit and global state:
 - Login / UGC endpoints - slower calls to a global state
 - Mutating endpoints - POP / regional state







Resiliency Techniques

Resiliency Techniques

- **Get Observability!**
 - Customized web logs with request metadata for modeling / triage



```

sub generate_bq_log {
    set resp.http.json generate bq_json = "{ "
    if(std.strlen(req.request) > 0, "%22" + "req_request" + "%22:%22" + json.escape(req.request) + "%22", "")
    if(std.strlen(req.service_id) > 0, ",%22" + "req_service_id" + "%22:%22" + json.escape(req.service_id) + "%22", "")
    if(std.strlen(req.backend.name) > 0, ",%22" + "req_backend" + "%22:%22" + json.escape(req.backend.name) + "%22", "")
    if(std.strlen(client.ip) > 0, ",%22" + "client_ip" + "%22:%22" + json.escape(client.ip) + "%22", "")
    ",%22" + "client_asn" + "%22:" + client.as.number
    ",%22" + "req_restarts" + "%22:" + req.restarts
    if(std.strlen(server.datacenter) > 0, ",%22" + "server_datacenter" + "%22:%22" + json.escape(server.datacenter) + "%22", "")
    if(std.strlen(fastly_info.state) > 0, ",%22" + "fastly_info_state" + "%22:%22" + json.escape(fastly_info.state) + "%22", "")
    if(std.strlen(resp.response) > 0, ",%22" + "resp_response" + "%22:%22" + json.escape(resp.response) + "%22", "")
    ",%22" + "resp_status" + "%22:" + resp.status
    ",%22" + "time_elapsed_msec" + "%22:" + time.elapsed.msec
    ",%22" + "time_end_msec" + "%22:" + time.end.msec
    ",%22" + "req_bytes_read" + "%22:" + req.bytes_read
    ",%22" + "resp_bytes_written" + "%22:" + resp.bytes_written
    ",%22" + "resp_body_bytes_written" + "%22:" + resp.body_bytes_written
    if(std.strlen(req.http.Host) > 0, ",%22" + "req_http_host" + "%22:%22" + json.escape(req.http.Host) + "%22", "")
    if(std.strlen(req.url) > 0, ",%22" + "req_url" + "%22:%22" + json.escape(req.url) + "%22", "")
    if(resp.http.bq-logging-skip-field:req_url_path,"",if(std.strlen(req.url.path) > 0, ",%22" + "req_url_path" + "%22:%22" + json.escape(req.url.path) + "%22", ""))
    if(resp.http.bq-logging-skip-field:req_url_qs,"",if(std.strlen(req.url.qs) > 0, ",%22" + "req_url_qs" + "%22:%22" + json.escape(req.url.qs) + "%22", ""))
    if(std.strlen(req.xid) > 0, ",%22" + "req_xid" + "%22:%22" + json.escape(req.xid) + "%22", "")
    if(std.strlen(req.method) > 0, ",%22" + "req_method" + "%22:%22" + json.escape(req.method) + "%22", "")
    if(std.strlen(substr(req.digest, 0, 16)) > 0, ",%22" + "req_cache_digest" + "%22:%22" + json.escape(substr(req.digest, 0, 16)) + "%22", "")
    if(std.strlen(req.http.CDN-Client-IP) > 0, ",%22" + "req_http_cdn_client_ip" + "%22:%22" + json.escape(req.http.CDN-Client-IP) + "%22", "")
    if(req.http.CDN-Client-IP && req.http.CDN-Client-IP ~ "^([0-9]{1,3}\.[0-9]{1,3}\.[0-9]{1,3}\.[0-9]{1,3})|([a-f0-9]{1,4})", ",%22" + "req_http_cdn_client_ip" + "%22", "")
    if(std.strlen(req.http.CDN-Org-Name) > 0, ",%22" + "req_http_cdn_org_name" + "%22:%22" + json.escape(req.http.CDN-Org-Name) + "%22", "")
    if(std.strlen(req.http.CDN-Proxy-Description) > 0, ",%22" + "req_http_cdn_proxy_description" + "%22:%22" + json.escape(req.http.CDN-Proxy-Description) + "%22", "")
    if(std.strlen(req.http.CDN-Proxy-Type) > 0, ",%22" + "req_http_cdn_proxy_type" + "%22:%22" + json.escape(req.http.CDN-Proxy-Type) + "%22", "")
    if(std.strlen(req.http.X-reddit-synthetic-code) > 0, ",%22" + "req_http_x_reddit_synthetic_code" + "%22:%22" + json.escape(req.http.X-reddit-synthetic-code) + "%22", "")
    if(std.strlen(req.http.X-reddit-synthetic-reason) > 0, ",%22" + "req_http_x_reddit_synthetic_reason" + "%22:%22" + json.escape(req.http.X-reddit-synthetic-reason) + "%22", "")
    if(std.strlen(req.http.X-reddit-pool-reason) > 0, ",%22" + "req_http_x_reddit_pool_reason" + "%22:%22" + json.escape(req.http.X-reddit-pool-reason) + "%22", "")
    if(std.strlen(req.http.X-reddit-block-reason) > 0, ",%22" + "req_http_x_reddit_block_reason" + "%22:%22" + json.escape(req.http.X-reddit-block-reason) + "%22", "")
    if(req.http.X-Reddit-Compression,"%22" + "req_http_x_reddit_compression" + "%22:" + json.escape(if(std.atoi(req.http.X-Reddit-Compression) != 1,"-1", "1")),"")
}

```



Resiliency Techniques

- **Get Observability!**
 - Customized web logs with request metadata for modeling / triage
 - Unique UUIDs on each request
 - Invest in tooling: Fastly terraform module for logging, tportctl



Resiliency Techniques

- Drop impossible requests (bad verbs, paths, synthetics to limit backend load)
- Session material validation at edge (JWT vs cookie)
- Resource pools
 - “Slowlane” and “woahlane” on trusted signal
 - Isolate degradation of service by request type / shape



Resiliency Techniques



- Make attacks more expensive for attackers:
 - **Tarpitting** - increase response time, requires attacker to have more concurrent port/request processing which might limit their size of attack
 - **Response Bloat** - best attacks in early Reddit were simple GET requests that cost nothing. Make them eat their network bandwidth at scale.


Resiliency Techniques

- **Cookie Trampolining**

- For brand new connections, add a cookie to a redirect response and see if client follows it
- Expected cookie not returned that should have had a request already?
- Good for blocking dumb scrapers / clients

|





Resiliency



Lessons Learned

- Want it done right, do/staff it yourself
- You don't need a WAF
- Expect cat and mouse
- Start with largest needles
- Want more insight? Buy votes





WAFS?!



WhAt AbOuT WAFS?!

- **Not a great fit** for Reddit - too many false positives, too many edge cases
 - HTML or SWE topics would trigger
- **Impacts performance, latency, and cost at our RPS scale**
- Better do it ourselves since we know our traffic patterns
- 🌶️ - Trade WAF for better appsec hygiene





Thanks

Made Possible with Contributions By:

Shadi Altarsha

Justin Ely

Clayton Gonsalves

Jason Harvey

Brian Landers

Udgeetha Mallampalli

Sotiris Nanopoulos

Sean Rees

Chris Schomaker

Hiram Silvey

